

# LAB 1: MNIST

**MNIST digit recognition using deep neural network**

한국과학기술원

시각지능연구실

박사과정 김태우

Enter the colab




google colab



 전체

 뉴스

 동영상

 이미지

 도서

 더보기

설정

도구

검색결과 약 3,120,000개 (0.38초)

colab.research.google.com ▾ 이 페이지 번역하기

## Google Colab

**Colab** notebooks execute code on **Google's** cloud servers, meaning you can leverage the power of **Google** hardware, including GPUs and TPUs, regardless of the ...

이 페이지를 20. 7. 29에 방문했습니다.

### Show notebooks in Drive

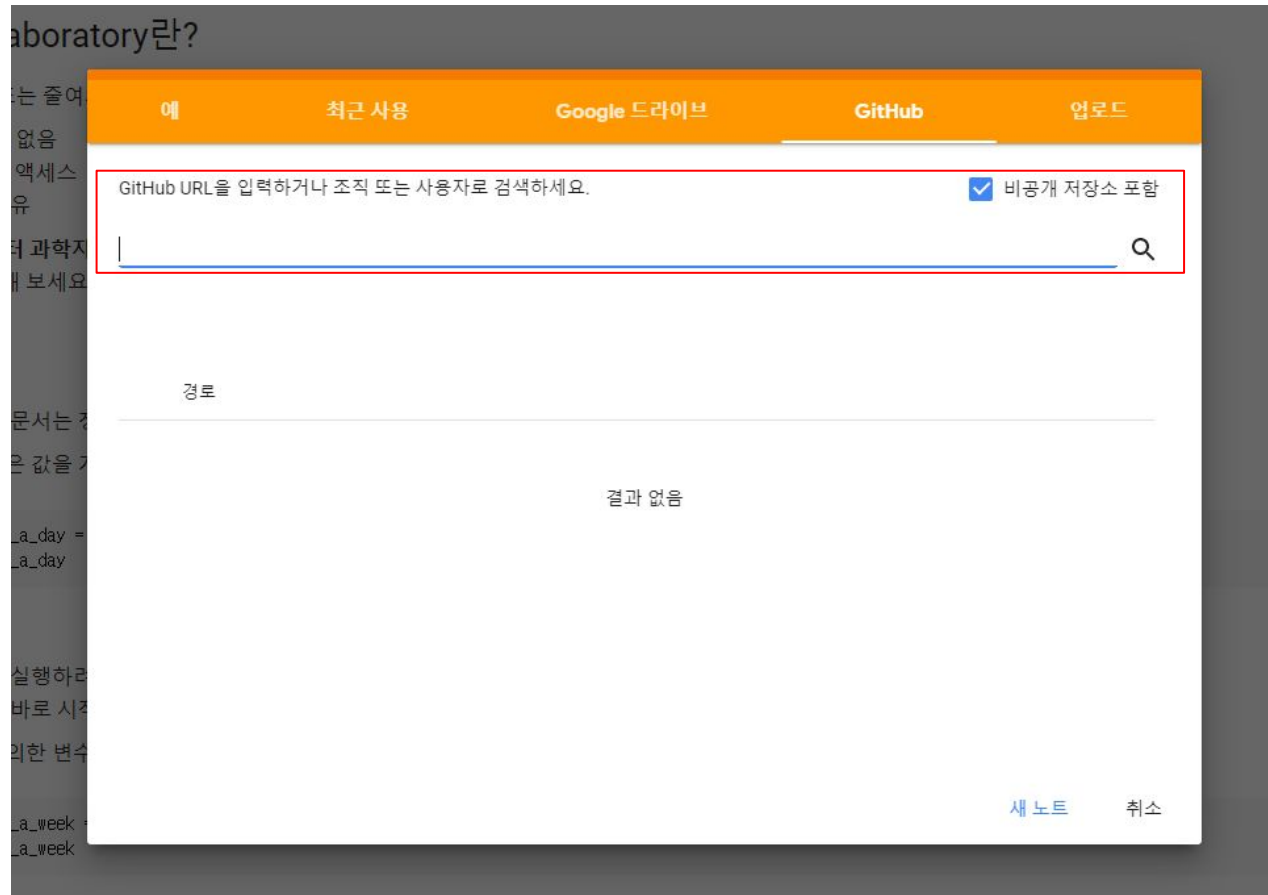
Sign in.

### Importing a library

To import a library that's not in Colaboratory by default, you ...

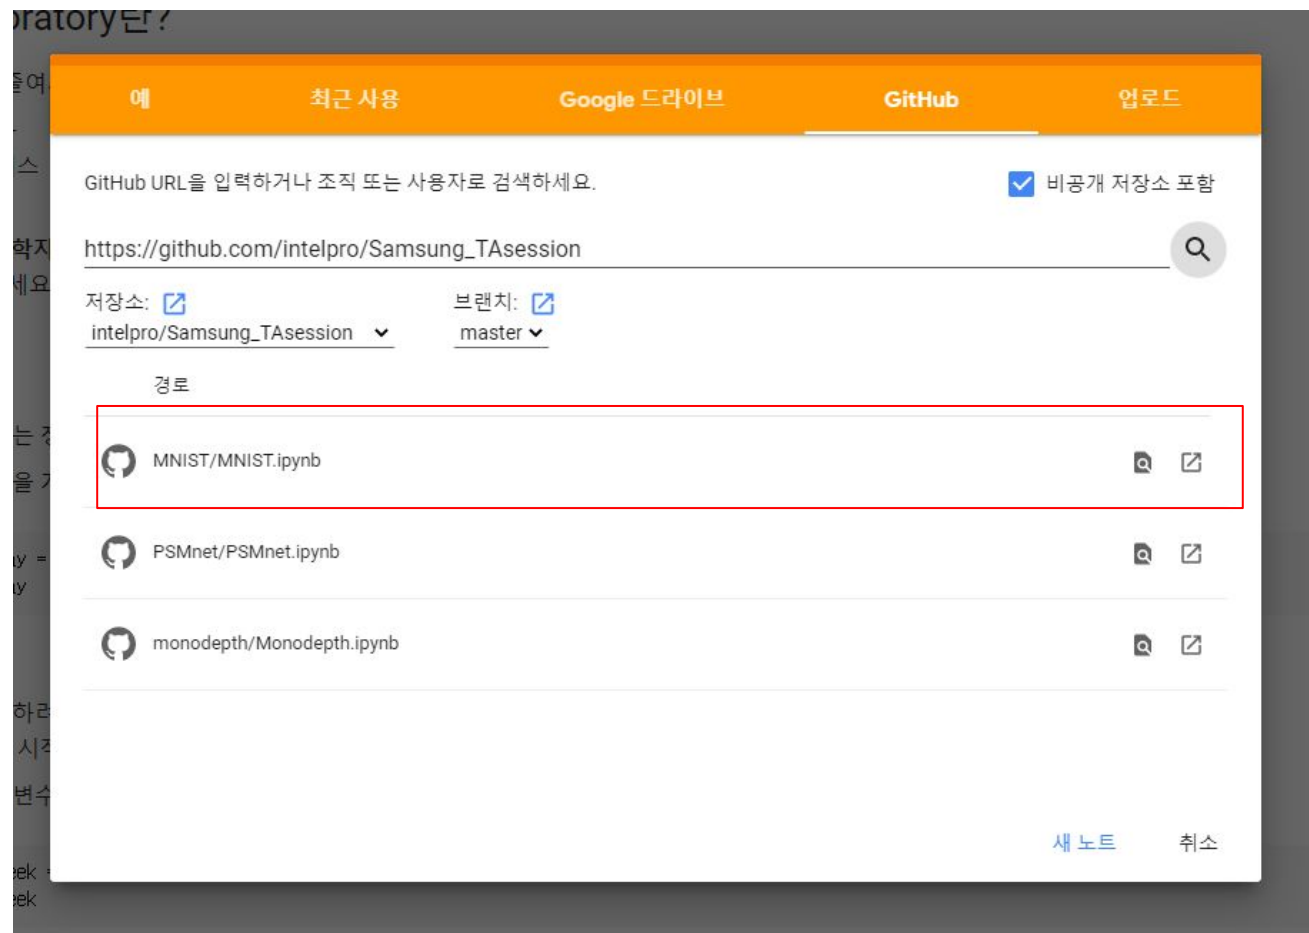
[google.com](#) 검색결과 더보기 »

Enter the colab + 아래의 주소를 입력한다.



[https://github.com/intelpro/Samsung\\_TAsession](https://github.com/intelpro/Samsung_TAsession)

## MNIST.ipynb 클릭



런타임 -> 런타임 유형 변경 -> 하드웨어  
가속기(GPU)

 \_MNIST.ipynb

파일 수정 보기 삽입 런타임 도구 도움말



## 노트 설정

하드웨어 가속기

GPU  

Colab를 최대한 활용하려면 필요하지 않은  
경우 GPU를 사용하지 않는 것이 좋습니다.

[자세히 알아보기](#)


☐ 이 노트를 저장할 때 코드 셀 출력 생략

취소

저장


## MNIST — Basic

- MNIST: Large hand written digit classification database
- Format
  - Input: 28 x 28 gray scale image
  - Output: 10 labels(0-9)
  - Centered on center of the mass




## Multi-layer perceptron

- Layer 1: Input layer. –  $(28^2 \times 1)$  dimension
- Layer 2: Hidden layer. – Multi-layer perceptron
- Layer 3: Output layer. – 10 labels(0-9)



## Multi-layer perceptron




$$[h'_1 \ h'_2 \ h'_3] = [x_1 \ x_2 \ x_3 \ \dots \ x_n] \cdot \begin{bmatrix} W_{11} & W_{12} & W_{13} \\ W_{21} & W_{22} & W_{23} \\ \vdots & \vdots & \vdots \\ W_{n1} & W_{n2} & W_{n3} \end{bmatrix}$$




## Multi-layer perceptron – dropout

- One of the biggest problems of deep learning is overfitting to training data.
- The dropout is a learning method using only a part of deep ne
- We can partially solve the overfitting problem through dropout method.



(a) Standard Neural Net



(b) After applying dropout.

## Pytorch MLP Basic – nn.Linear

**CLASS** `torch.nn.Linear(in_features: int, out_features: int, bias: bool = True)`

[SOURCE]

Applies a linear transformation to the incoming data:  $y = xA^T + b$

### Parameters

- **in\_features** – size of each input sample
- **out\_features** – size of each output sample
- **bias** – If set to `False`, the layer will not learn an additive bias. Default: `True`

### Shape:

- Input:  $(N, *, H_{in})$  where  $*$  means any number of additional dimensions and  $H_{in} = \text{in\_features}$
- Output:  $(N, *, H_{out})$  where all but the last dimension are the same shape as the input and  $H_{out} = \text{out\_features}$ .

### Variables

- **~Linear.weight** – the learnable weights of the module of shape  $(\text{out\_features}, \text{in\_features})$ . The values are initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ , where  $k = \frac{1}{\text{in\_features}}$
- **~Linear.bias** – the learnable bias of the module of shape  $(\text{out\_features})$ . If `bias` is `True`, the values are initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where  $k = \frac{1}{\text{in\_features}}$

<https://pytorch.org/docs/master/generated/torch.nn.Linear.html>

## Pytorch MLP Basic – nn.Dropout

**CLASS** `torch.nn.Dropout(p: float = 0.5, inplace: bool = False)`

[\[SOURCE\]](#)

During training, randomly zeroes some of the elements of the input tensor with probability `p` using samples from a Bernoulli distribution. Each channel will be zeroed out independently on every forward call.

This has proven to be an effective technique for regularization and preventing the co-adaptation of neurons as described in the paper [Improving neural networks by preventing co-adaptation of feature detectors](#).

Furthermore, the outputs are scaled by a factor of  $\frac{1}{1-p}$  during training. This means that during evaluation the module simply computes an identity function.

### Parameters

- **p** – probability of an element to be zeroed. Default: 0.5
- **inplace** – If set to `True`, will do this operation in-place. Default: `False`

### Shape:

- Input:  $(*)$ . Input can be of any shape
- Output:  $(*)$ . Output is of the same shape as input

<https://pytorch.org/docs/master/generated/torch.nn.Dropout.html>

## MNIST tutorial

### Dataset preparation

```
batch_size = 32

kwargs = {'num_workers': 1, 'pin_memory': True} if cuda else {}

train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=True, download=True,
                   transform=transforms.Compose([
                       transforms.ToTensor(),
                       transforms.Normalize((0.1307,), (0.3081,))
                   ])),
    batch_size=batch_size, shuffle=True, **kwargs)

validation_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=False, transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))
    ])),
    batch_size=batch_size, shuffle=False, **kwargs)
```

- Load MNIST dataset using data loader
- Input size: 28 x 28 gray scale image
- Output: classes of ("0", .. "9") for each training digit

## MNIST tutorial

### Network definition class

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(28*28, 50)
        self.fc1_drop = nn.Dropout(0.2)
        self.fc2 = nn.Linear(50, 50)
        self.fc2_drop = nn.Dropout(0.2)
        self.fc3 = nn.Linear(50, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = F.relu(self.fc1(x))
        x = self.fc1_drop(x)
        x = F.relu(self.fc2(x))
        x = self.fc2_drop(x)
        return F.log_softmax(self.fc3(x))
```

Network definition

Forward function of  
input data

— Activation  
function

## MNIST tutorial

### Training

```
def train(epoch, log_interval=100):  
    model.train()  
    for batch_idx, (data, target) in enumerate(train_loader):  
        if cuda:  
            data, target = data.cuda(), target.cuda()  
            data, target = Variable(data), Variable(target)  
            optimizer.zero_grad()  
            output = model(data) → Get output  
            loss = F.nll_loss(output, target) → Calculate loss  
            loss.backward()  
            optimizer.step() } → Backpropagation using optimizer  
        if batch_idx % log_interval == 0:  
            print('Train Epoch: {} [{}/{} ({:.0f}%)] \t Loss: {:.6f}'.format(  
                epoch, batch_idx * len(data), len(train_loader.dataset),  
                100. * batch_idx / len(train_loader), loss.data[0])) → Print training log
```



## MNIST tutorial

### Validation

```
def validate(loss_vector, accuracy_vector):
    model.eval()
    val_loss, correct = 0, 0
    for data, target in validation_loader:
        if cuda:
            data, target = data.cuda(), target.cuda()
        data, target = Variable(data, volatile=True), Variable(target)
        output = model(data)
        val_loss += F.nll_loss(output, target).data[0]
        pred = output.data.max(1)[1] # get the index of the max log-probability
        correct += pred.eq(target.data).cpu().sum()

    val_loss /= len(validation_loader)
    loss_vector.append(val_loss)

    accuracy = 100. * correct / len(validation_loader.dataset)
    accuracy_vector.append(accuracy)

    print('\nValidation set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
        val_loss, correct, len(validation_loader.dataset), accuracy))
```

Calculate validation loss

Calculate accuracy of validation set



Print log of accuracy of validation set

## MNIST tutorial

```
plt.figure(figsize=(5,3))
plt.plot(np.arange(1,epochs+1), lossv)
plt.title('validation loss')

plt.figure(figsize=(5,3))
plt.plot(np.arange(1,epochs+1), accv)
plt.title('validation accuracy');
```

**Plot the loss and accuracy graph**  
We can visualize the performance and the loss of our network in validation sets.





## MNIST with Simple CNN

### CONV2D

```
CLASS torch.nn.Conv2d(in_channels: int, out_channels: int, kernel_size: Union[T,  
    Tuple[T, T]], stride: Union[T, Tuple[T, T]] = 1, padding: Union[T, Tuple[T, T]] =  
    0, dilation: Union[T, Tuple[T, T]] = 1, groups: int = 1, bias: bool = True,  
    padding_mode: str = 'zeros') [SOURCE]
```

#### Parameters

- **in\_channels** (*int*) – Number of channels in the input image
- **out\_channels** (*int*) – Number of channels produced by the convolution
- **kernel\_size** (*int or tuple*) – Size of the convolving kernel
- **stride** (*int or tuple, optional*) – Stride of the convolution. Default: 1
- **padding** (*int or tuple, optional*) – Zero-padding added to both sides of the input. Default: 0
- **padding\_mode** (*string, optional*) – 'zeros', 'reflect', 'replicate' or 'circular'. Default: 'zeros'
- **dilation** (*int or tuple, optional*) – Spacing between kernel elements. Default: 1
- **groups** (*int, optional*) – Number of blocked connections from input channels to output channels. Default: 1
- **bias** (*bool, optional*) – If `True`, adds a learnable bias to the output. Default: `True`

## MNIST with Simple CNN

Shape:

- Input:  $(N, C_{in}, H_{in}, W_{in})$
- Output:  $(N, C_{out}, H_{out}, W_{out})$  where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel\_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel\_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

Examples

```
>>> # With square kernels and equal stride
>>> m = nn.Conv2d(16, 33, 3, stride=2)
>>> # non-square kernels and unequal stride and with padding
>>> m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2))
>>> # non-square kernels and unequal stride and with padding and dilation
>>> m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2), dilation=(3, 1))
>>> input = torch.randn(20, 16, 50, 100)
>>> output = m(input)
```

## MNIST with Simple CNN

### MAXPOOL2D

```
CLASS torch.nn.MaxPool2d(kernel_size: Union[T, Tuple[T, ...]], stride:
Optional[Union[T, Tuple[T, ...]]] = None, padding: Union[T, Tuple[T, ...]] = 0,
dilation: Union[T, Tuple[T, ...]] = 1, return_indices: bool = False, ceil_mode:
bool = False) [SOURCE]
```

Applies a 2D max pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C, H, W)$ , output  $(N, C, H_{out}, W_{out})$  and `kernel_size`  $(kH, kW)$  can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} input(N_i, C_j, stride[0] \times h + m, stride[1] \times w + n)$$

If `padding` is non-zero, then the input is implicitly zero-padded on both sides for `padding` number of points. `dilation` controls the spacing between the kernel points. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.

The parameters `kernel_size`, `stride`, `padding`, `dilation` can either be:

- a single `int` – in which case the same value is used for the height and width dimension
- a `tuple` of two ints – in which case, the first `int` is used for the height dimension, and the second `int` for the width dimension

## MNIST with Simple CNN

Shape:

- Input:  $(N, C, H_{in}, W_{in})$
- Output:  $(N, C, H_{out}, W_{out})$ , where

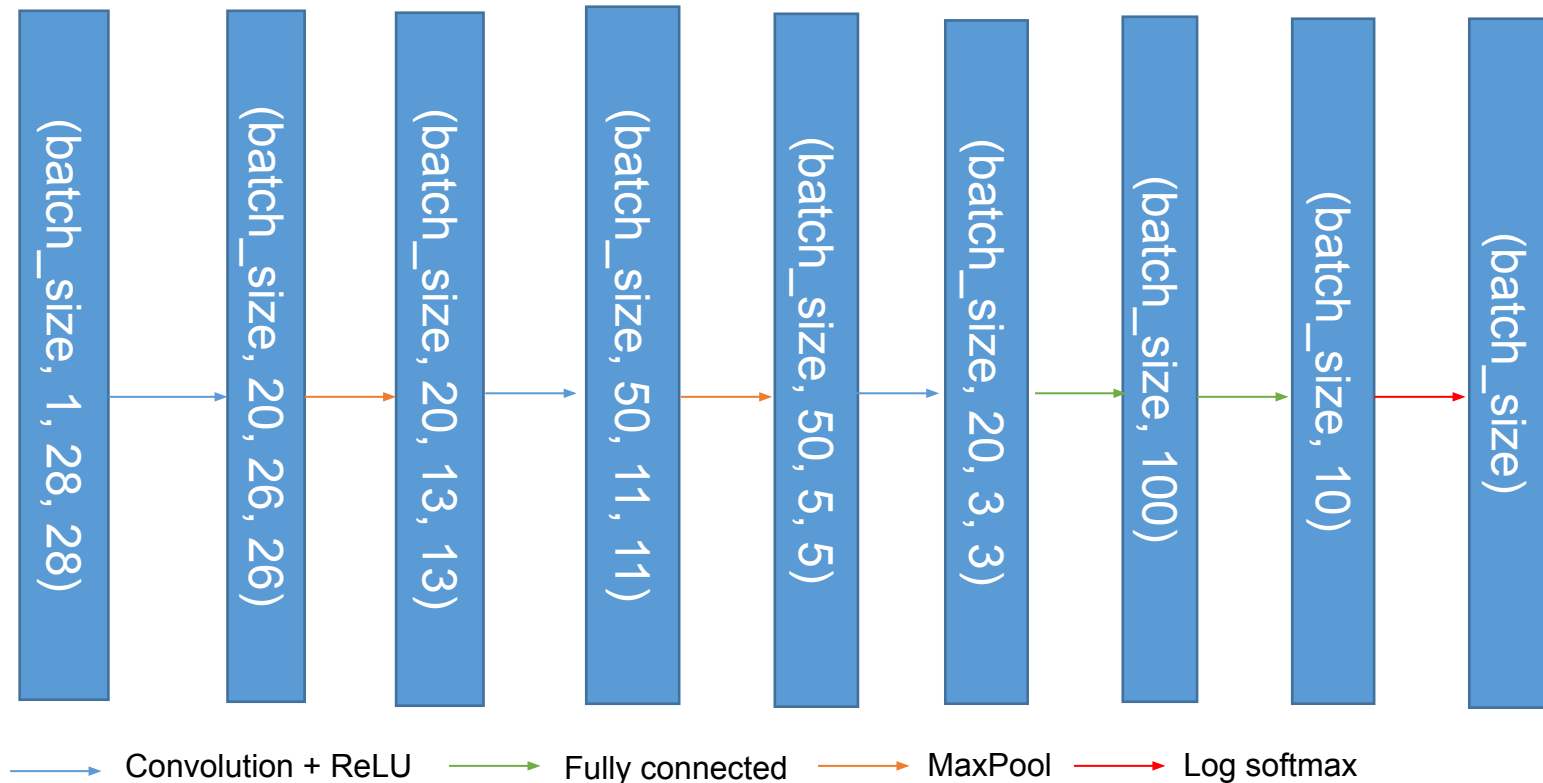
$$H_{out} = \left\lfloor \frac{H_{in} + 2 * padding[0] - dilation[0] \times (kernel\_size[0] - 1) - 1}{stride[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 * padding[1] - dilation[1] \times (kernel\_size[1] - 1) - 1}{stride[1]} + 1 \right\rfloor$$

Examples:

```
>>> # pool of square window of size=3, stride=2
>>> m = nn.MaxPool2d(3, stride=2)
>>> # pool of non-square window
>>> m = nn.MaxPool2d((3, 2), stride=(2, 1))
>>> input = torch.randn(20, 16, 50, 32)
>>> output = m(input)
```

## MNIST with Simple CNN



### Note:

Learning with simple CNN의 Notimplemented부분을 채워넣어 주세요

Caution: when you use construct CNN, do not use stride or padding size, Just use proper kernel size to match dimension

## MNIST with VGG like-CNN

layer	output dimension	layer	output dimension
0	(batch_size, 1, 28, 28)	13 - Linear1	(batch_size,100)
1 - Conv + ReLU	(batch_size, 64, 30, 30)	14 - Dropout	(batch_size,100)
2 - max pooling 2D(2x2)	(batch_size, 64, 15, 15)	15 - Linear2	(batch_size,100)
3 - Conv + ReLU	(batch_size, 128, 17, 17)	16 - Dropout	(batch_size,10)
4 - max pooling 2D(2x2)	(batch_size, 128, 8, 8)	17 - Linear3	(batch_size)
5 - Conv + Conv + ReLU	(batch_size, 256, 10, 10)	<b>Note:</b> Learning with VGG like CNN의 Notimplemented를 채워주세요 주황색으로 표시된 부분이 구현해야 할 부분입니다.	
7 - max pooling 2D(2x2)	(batch_size, 256, 5, 5)		
8 - Conv + Conv + ReLU	(batch_size, 512, 5, 5)		
9 - max pooling 2D(2x2)	(batch_size, 512, 2, 2)		
10 - Conv + Conv + ReLU	(batch_size, 512, 2, 2)		
11 - max pool 2D(2x2)	(batch_size, 512, 1, 1)		
12 - stretch tensor	(batch_size, 512)		