

LAB 1: MNIST

MNIST digit recognition using deep neural network


한국과학기술원

시각지능연구실

박사과정 김태우


MNIST — Basic

- MNIST: Large hand written digit classification database
- Format
 - Input: 28 x 28 gray scale image
 - Output: 10 labels(0-9)
 - Centered on center of the mass




Multi-layer perceptron

- Layer 1: Input layer. – $(28^2 \times 1)$ dimension
- Layer 2: Hidden layer. – Multi-layer perceptron
- Layer 3: Output layer. – 10 labels(0-9)




Multi-layer perceptron




$$[h'_1 \ h'_2 \ h'_3] = [x_1 \ x_2 \ x_3 \ \dots \ x_n] \cdot \begin{bmatrix} W_{11} & W_{12} & W_{13} \\ W_{21} & W_{22} & W_{23} \\ \vdots & \vdots & \vdots \\ W_{n1} & W_{n2} & W_{n3} \end{bmatrix}$$

Multi-layer perceptron – dropout

- One of the biggest problems of deep learning is overfitting to training data.
- The dropout is a learning method using only a part of deep ne
- We can partially solve the overfitting problem through dropout method.



(a) Standard Neural Net



(b) After applying dropout.

Pytorch MLP Basic – nn.Linear

```
CLASS torch.nn.Linear(in_features: int, out_features: int, bias: bool = True)
```

[\[SOURCE\]](#)

Applies a linear transformation to the incoming data: $y = xA^T + b$

Parameters

- **in_features** – size of each input sample
- **out_features** – size of each output sample
- **bias** – If set to `False`, the layer will not learn an additive bias. Default: `True`

Shape:

- Input: $(N, *, H_{in})$ where $*$ means any number of additional dimensions and $H_{in} = \text{in_features}$
- Output: $(N, *, H_{out})$ where all but the last dimension are the same shape as the input and $H_{out} = \text{out_features}$.

Variables

- **~Linear.weight** – the learnable weights of the module of shape $(\text{out_features}, \text{in_features})$. The values are initialized from $\mathcal{U}(-\sqrt{k}, \sqrt{k})$, where $k = \frac{1}{\text{in_features}}$
- **~Linear.bias** – the learnable bias of the module of shape (out_features) . If `bias` is `True`, the values are initialized from $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ where $k = \frac{1}{\text{in_features}}$

Pytorch MLP Basic – nn.Dropout

CLASS `torch.nn.Dropout(p: float = 0.5, inplace: bool = False)`

[\[SOURCE\]](#)

During training, randomly zeroes some of the elements of the input tensor with probability `p` using samples from a Bernoulli distribution. Each channel will be zeroed out independently on every forward call.

This has proven to be an effective technique for regularization and preventing the co-adaptation of neurons as described in the paper [Improving neural networks by preventing co-adaptation of feature detectors](#).

Furthermore, the outputs are scaled by a factor of $\frac{1}{1-p}$ during training. This means that during evaluation the module simply computes an identity function.

Parameters

- **p** – probability of an element to be zeroed. Default: 0.5
- **inplace** – If set to `True`, will do this operation in-place. Default: `False`

Shape:

- Input: $(*)$. Input can be of any shape
- Output: $(*)$. Output is of the same shape as input

MNIST tutorial

Dataset preparation

```
batch_size = 32

kwargs = {'num_workers': 1, 'pin_memory': True} if cuda else {}

train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=True, download=True,
                   transform=transforms.Compose([
                       transforms.ToTensor(),
                       transforms.Normalize((0.1307,), (0.3081,))
                   ])),
    batch_size=batch_size, shuffle=True, **kwargs)

validation_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=False, transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))
    ])),
    batch_size=batch_size, shuffle=False, **kwargs)
```

- Load MNIST dataset using data loader
- Input size: 28 x 28 gray scale image
- Output: classes of ("0", .. "9") for each training digit

MNIST tutorial

Network definition class

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(28*28, 50)
        self.fc1_drop = nn.Dropout(0.2)
        self.fc2 = nn.Linear(50, 50)
        self.fc2_drop = nn.Dropout(0.2)
        self.fc3 = nn.Linear(50, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = F.relu(self.fc1(x))
        x = self.fc1_drop(x)
        x = F.relu(self.fc2(x))
        x = self.fc2_drop(x)
        return F.log_softmax(self.fc3(x))
```

Network definition

Feed forward of
input data

Activation
function

MNIST tutorial

Training

```
def train(epoch, log_interval=100):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        if cuda:
            data, target = data.cuda(), target.cuda()
            data, target = Variable(data), Variable(target)
            optimizer.zero_grad()
            output = model(data) → Get output
            loss = F.nll_loss(output, target) → Calculate loss
            loss.backward()
            optimizer.step() } → Backpropagation using optimizer
        if batch_idx % log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)] \t Loss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.data[0])) → Print training log
```

MNIST tutorial

Validation

```
def validate(loss_vector, accuracy_vector):
    model.eval()
    val_loss, correct = 0, 0
    for data, target in validation_loader:
        if cuda:
            data, target = data.cuda(), target.cuda()
        data, target = Variable(data, volatile=True), Variable(target)
        output = model(data)
        val_loss += F.nll_loss(output, target).data[0]
        pred = output.data.max(1)[1] # get the index of the max log-probability
        correct += pred.eq(target.data).cpu().sum()

    val_loss /= len(validation_loader)
    loss_vector.append(val_loss)

    accuracy = 100. * correct / len(validation_loader.dataset)
    accuracy_vector.append(accuracy)

    print('\nValidation set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
        val_loss, correct, len(validation_loader.dataset), accuracy))
```

Calculate validation loss

Calculate accuracy of validation set



Print log of accuracy of validation set

MNIST tutorial

```
plt.figure(figsize=(5,3))
plt.plot(np.arange(1,epochs+1), lossv)
plt.title('validation loss')

plt.figure(figsize=(5,3))
plt.plot(np.arange(1,epochs+1), accv)
plt.title('validation accuracy');
```

Plot the loss and accuracy graph
We can visualize the performance and the loss of our network in validation sets.



MNIST with Simple CNN

CONV2D

```
CLASS torch.nn.Conv2d(in_channels: int, out_channels: int, kernel_size: Union[T,  
    Tuple[T, T]], stride: Union[T, Tuple[T, T]] = 1, padding: Union[T, Tuple[T, T]] =  
    0, dilation: Union[T, Tuple[T, T]] = 1, groups: int = 1, bias: bool = True,  
    padding_mode: str = 'zeros') [SOURCE]
```

Parameters

- **in_channels** (*int*) – Number of channels in the input image
- **out_channels** (*int*) – Number of channels produced by the convolution
- **kernel_size** (*int or tuple*) – Size of the convolving kernel
- **stride** (*int or tuple, optional*) – Stride of the convolution. Default: 1
- **padding** (*int or tuple, optional*) – Zero-padding added to both sides of the input. Default: 0
- **padding_mode** (*string, optional*) – 'zeros', 'reflect', 'replicate' or 'circular'. Default: 'zeros'
- **dilation** (*int or tuple, optional*) – Spacing between kernel elements. Default: 1
- **groups** (*int, optional*) – Number of blocked connections from input channels to output channels. Default: 1
- **bias** (*bool, optional*) – If `True`, adds a learnable bias to the output. Default: `True`

MNIST with Simple CNN

Shape:

- Input: $(N, C_{in}, H_{in}, W_{in})$
- Output: $(N, C_{out}, H_{out}, W_{out})$ where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

Examples

```
>>> # With square kernels and equal stride
>>> m = nn.Conv2d(16, 33, 3, stride=2)
>>> # non-square kernels and unequal stride and with padding
>>> m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2))
>>> # non-square kernels and unequal stride and with padding and dilation
>>> m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2), dilation=(3, 1))
>>> input = torch.randn(20, 16, 50, 100)
>>> output = m(input)
```


MNIST with Simple CNN

MAXPOOL2D

```
CLASS torch.nn.MaxPool2d(kernel_size: Union[T, Tuple[T, ...]], stride:
Optional[Union[T, Tuple[T, ...]]] = None, padding: Union[T, Tuple[T, ...]] = 0,
dilation: Union[T, Tuple[T, ...]] = 1, return_indices: bool = False, ceil_mode:
bool = False)
```

[SOURCE]

Applies a 2D max pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C, H, W) , output (N, C, H_{out}, W_{out}) and `kernel_size` (kH, kW) can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} input(N_i, C_j, stride[0] \times h + m, stride[1] \times w + n)$$

If `padding` is non-zero, then the input is implicitly zero-padded on both sides for `padding` number of points. `dilation` controls the spacing between the kernel points. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.

The parameters `kernel_size`, `stride`, `padding`, `dilation` can either be:

- a single `int` – in which case the same value is used for the height and width dimension
- a `tuple` of two ints – in which case, the first `int` is used for the height dimension, and the second `int` for the width dimension

MNIST with Simple CNN

Shape:

- Input: (N, C, H_{in}, W_{in})
- Output: (N, C, H_{out}, W_{out}) , where

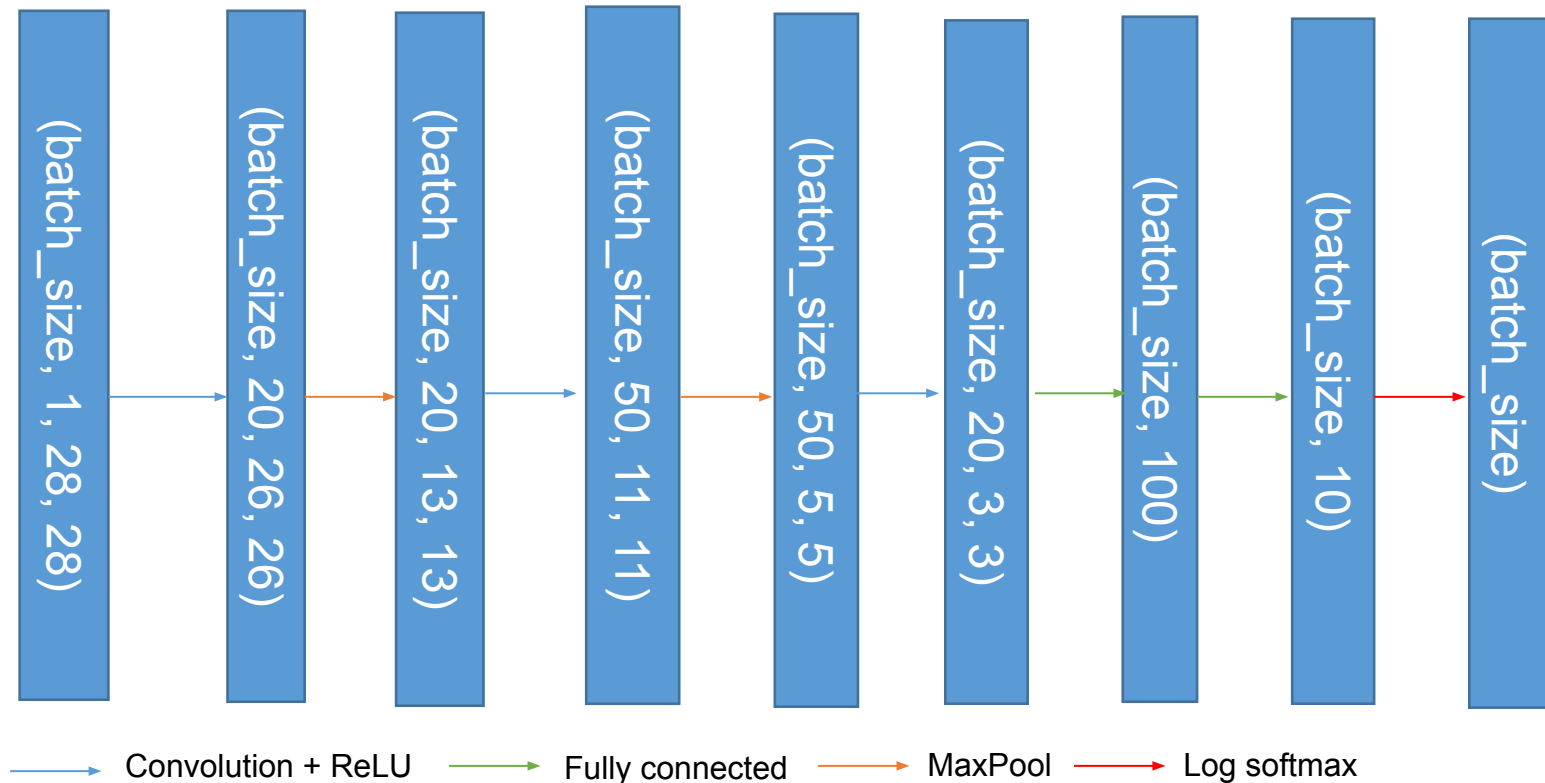
$$H_{out} = \left\lfloor \frac{H_{in} + 2 * padding[0] - dilation[0] \times (kernel_size[0] - 1) - 1}{stride[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 * padding[1] - dilation[1] \times (kernel_size[1] - 1) - 1}{stride[1]} + 1 \right\rfloor$$

Examples:

```
>>> # pool of square window of size=3, stride=2
>>> m = nn.MaxPool2d(3, stride=2)
>>> # pool of non-square window
>>> m = nn.MaxPool2d((3, 2), stride=(2, 1))
>>> input = torch.randn(20, 16, 50, 32)
>>> output = m(input)
```

MNIST with Simple CNN



Caution: when you use construct CNN, do not use stride or padding size, Just use proper kernel size to match dimension


MNIST with VGG like-CNN

layer	output dimension	layer	output dimension
0	(batch_size, 1, 28, 28)	13 - Linear1	(batch_size,100)
1 - Conv + ReLU	(batch_size, 64, 30, 30)	14 - Droput	(batch_size,100)
2 - max pooling 2D(2x2)	(batch_size, 64, 15, 15)	15 - Linear2	(batch_size,100)
3 - Conv + ReLU	(batch_size, 128, 17, 17)	16 - Dropout	(batch_size,10)
4 - max pooling 2D(2x2)	(batch_size, 128, 8, 8)	17 - Linear3	(batch_size)
5 - Conv + Conv + ReLU	(batch_size, 256, 10, 10)		
7 - max pooling 2D(2x2)	(batch_size, 256, 5, 5)		
8 - Conv + Conv + ReLU	(batch_size, 512, 5, 5)		
9 - max pooling 2D(2x2)	(batch_size, 512, 2, 2)		
10 - Conv + Conv + ReLU	(batch_size, 512, 2, 2)		
11 - max pool 2D(2x2)	(batch_size, 512, 1, 1)		
12 - stretch tensor	(batch_size, 512)		

Enter the colab



google colab



검색결과 약 3,120,000개 (0.38초)

colab.research.google.com ▾ 이 페이지 번역하기

Google Colab

Colab notebooks execute code on **Google's** cloud servers, meaning you can leverage the power of **Google** hardware, including GPUs and TPUs, regardless of the ...

이 페이지를 20. 7. 29에 방문했습니다.

Show notebooks in Drive

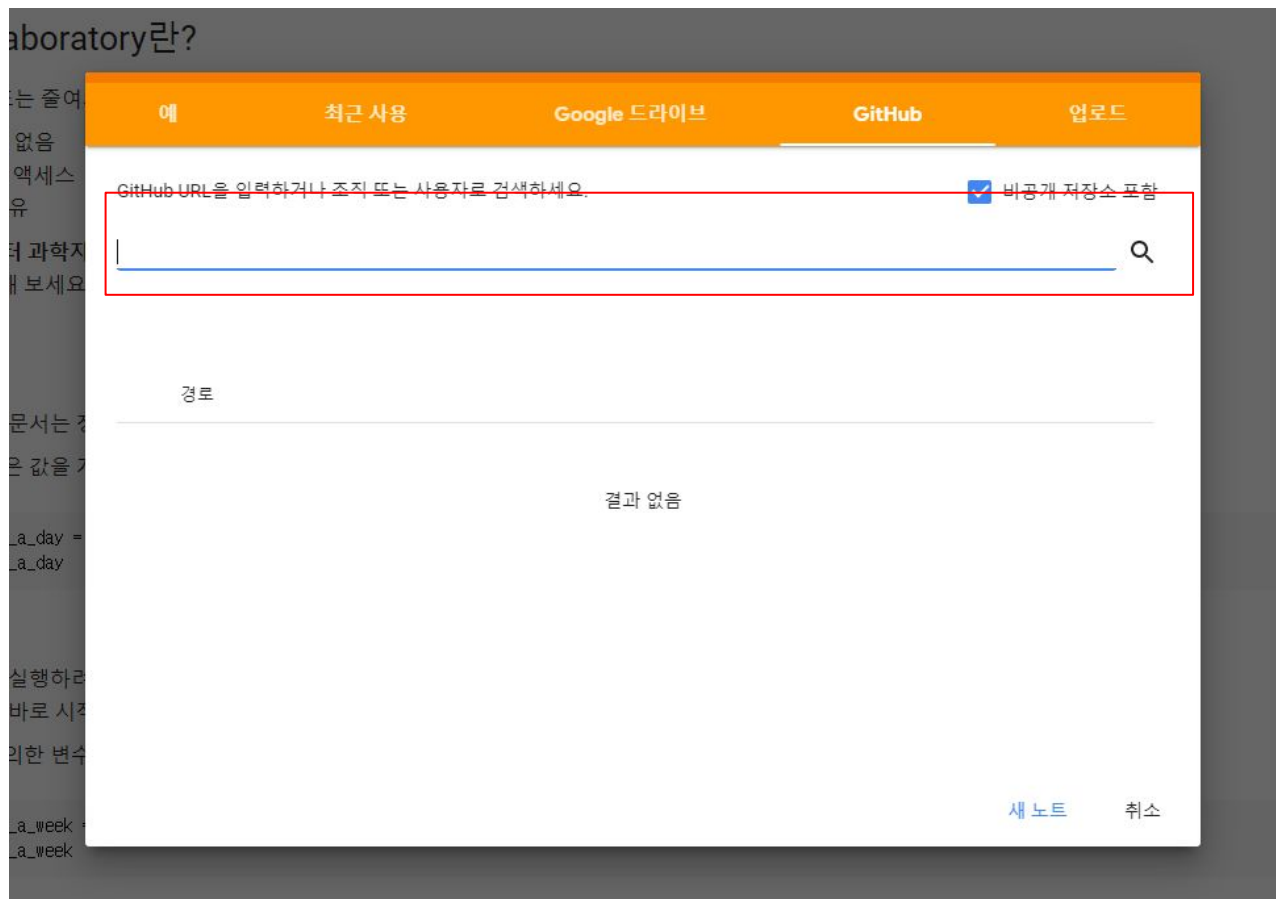
Sign in.

Importing a library

To import a library that's not in Colaboratory by default, you ...

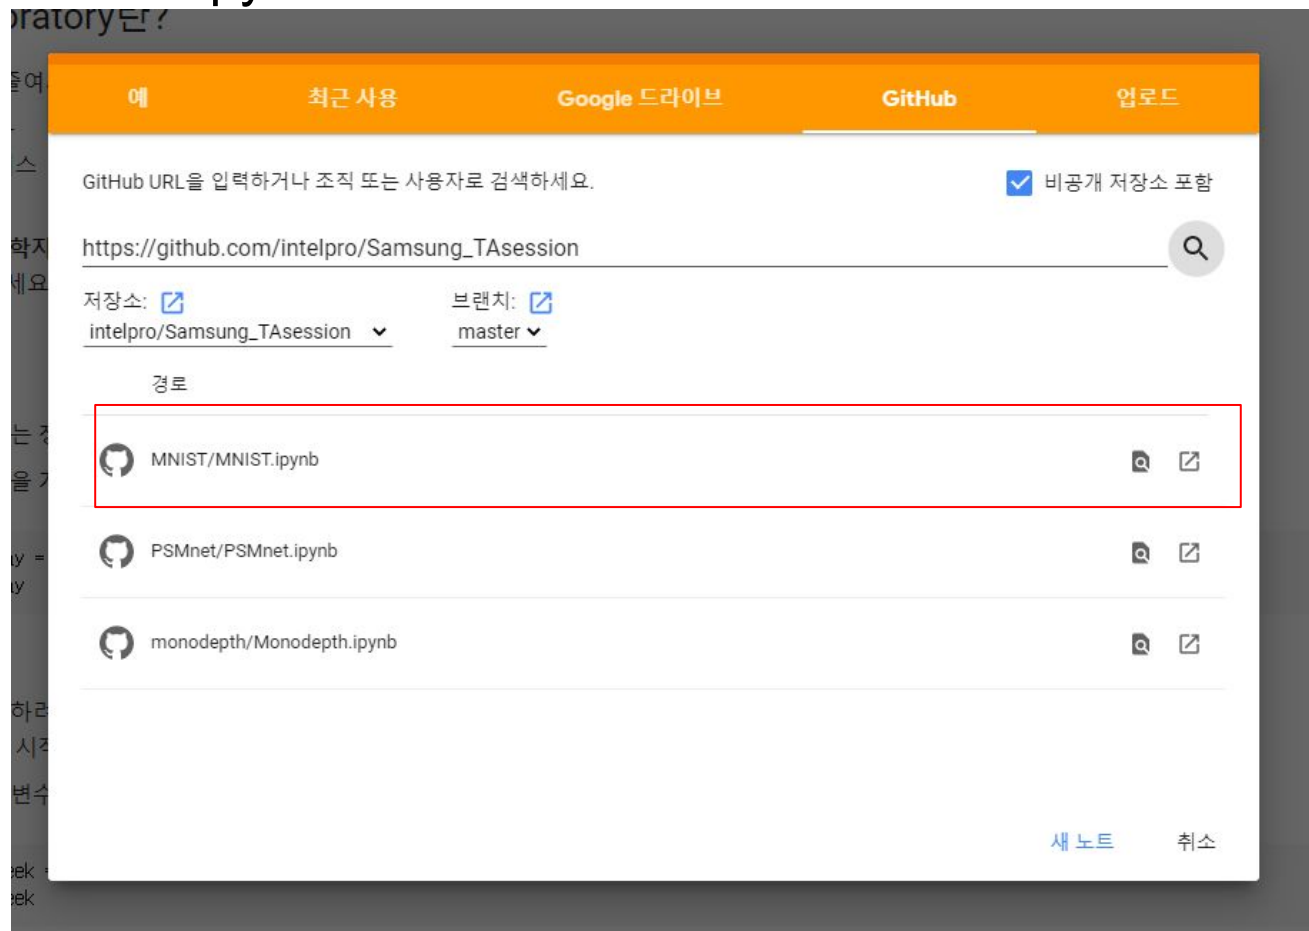
[google.com](#) 검색결과 더보기 »

Enter the colab + 아래의 주소를 입력한다.

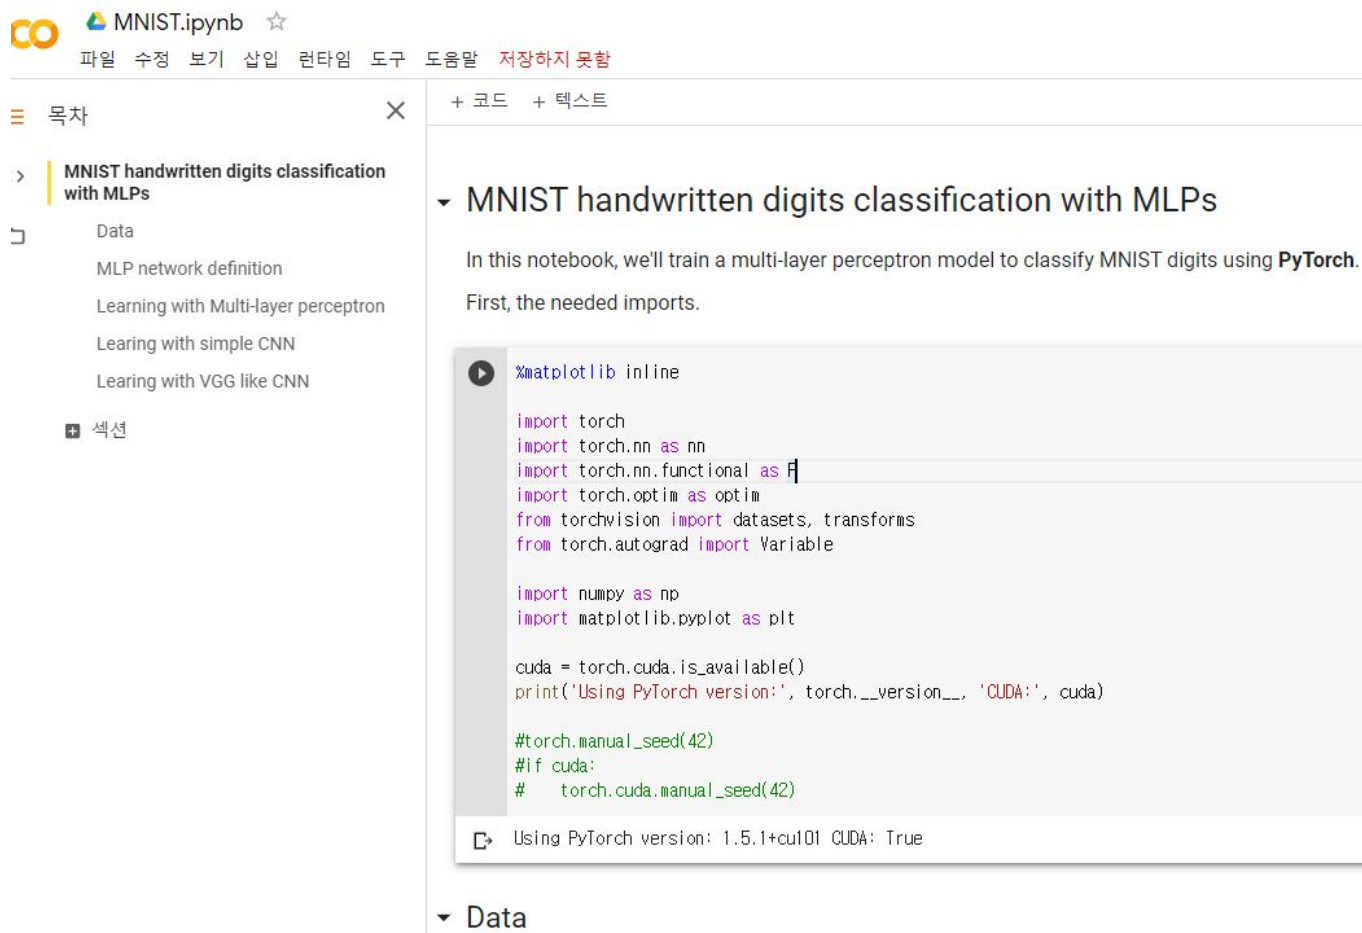


https://github.com/intelpro/Samsung_TAsession

MNIST.ipynb 클릭



런타임 -> 런타임 유형 변경 -> 하드웨어 가속기(GPU)



The screenshot shows a Google Colab notebook interface. At the top, the notebook is titled 'MNIST.ipynb' with a star icon. Below the title, there are tabs for '파일' (File), '수정' (Edit), '보기' (View), '삽입' (Insert), '런타임' (Runtime), '도구' (Tools), '도움말' (Help), and '저장하지 못함' (Can't save). The '런타임' tab is active, and a dropdown menu is open, showing options to change the runtime type. The current runtime type is 'Python 3 (TensorFlow 1.15.0)' and the hardware accelerator is 'None'. The dropdown menu shows 'Python 3 (TensorFlow 1.15.0)' and 'Python 3 (TensorFlow 1.15.0) (GPU)' as options. The 'Python 3 (TensorFlow 1.15.0) (GPU)' option is highlighted. Below the dropdown, there is a button that says '변경' (Change).

On the left sidebar, the notebook's table of contents is visible, showing sections like 'Data', 'MLP network definition', 'Learning with Multi-Layer perceptron', 'Learning with simple CNN', and 'Learning with VGG like CNN'. The 'Data' section is currently selected.

The main content area of the notebook shows the title 'MNIST handwritten digits classification with MLPs'. Below the title, there is a text block that says 'In this notebook, we'll train a multi-layer perceptron model to classify MNIST digits using **PyTorch**. First, the needed imports.'

Below the text block, there is a code cell with the following code:

```
%matplotlib inline

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
from torch.autograd import Variable

import numpy as np
import matplotlib.pyplot as plt

cuda = torch.cuda.is_available()
print('Using PyTorch version:', torch.__version__, 'CUDA:', cuda)

#torch.manual_seed(42)
#if cuda:
#    torch.cuda.manual_seed(42)
```

Below the code cell, there is a status bar that says 'Using PyTorch version: 1.5.1+cu101 CUDA: True'.